

Question-1. Download dataset

Question-2. Load the dataset

```
path = '/content/drive/MyDrive/Colab Notebooks/Churn_Modelling.csv'
```

```
import pandas as pd
```

```
data = pd.read_csv('/content/drive/MyDrive/Colab  
Notebooks/Churn_Modelling.csv')
```

```
data.head
```

```
<bound method NDFrame.head of  
CreditScore Geography Gender Age \  
0          1      15634602   Hargrave      619      France  Female  
42  
1          2      15647311      Hill      608      Spain  Female  
41  
2          3      15619304      Onio      502      France  Female  
42  
3          4      15701354      Boni      699      France  Female  
39  
4          5      15737888  Mitchell      850      Spain  Female  
43  
...      ...      ...      ...      ...      ...      ...  
...  
9995      9996      15606229  Obijiaku      771      France   Male  
39  
9996      9997      15569892  Johnstone      516      France   Male  
35  
9997      9998      15584532      Liu      709      France  Female  
36  
9998      9999      15682355  Sabbatini      772      Germany   Male  
42  
9999      10000      15628319   Walker      792      France  Female  
28
```

```
Tenure      Balance  NumOfProducts  HasCrCard  IsActiveMember  \  
0          2         0.00              1           1              1  
1          1      83807.86              1           0              1  
2          8     159660.80              3           1              0  
3          1         0.00              2           0              0  
4          2     125510.82              1           1              1  
...      ...      ...      ...      ...      ...  
9995        5         0.00              2           1              0  
9996       10      57369.61              1           1              1  
9997        7         0.00              1           0              1  
9998        3      75075.31              2           1              0  
9999        4     130142.79              1           1              0
```

	EstimatedSalary	Exited
0	101348.88	1
1	112542.58	0
2	113931.57	1
3	93826.63	0
4	79084.10	0
...
9995	96270.64	0
9996	101699.77	0
9997	42085.58	1
9998	92888.52	1
9999	38190.78	0


[10000 rows x 14 columns]>

Question-3. Perform Below Visualizations.

3.1 univariate analysis


```
sns.displot(data['CreditScore'])
```

```
<seaborn.axisgrid.FacetGrid at 0x7ff47b124f50>
```




```
sns.histplot(data['CreditScore'])
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7ff47aec4a50>
```



```
sns.boxplot(x = data['CreditScore'])
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7ff4780406d0>
```




```
sns.distplot(data['Age'])
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/distributions.py:2619:  
FutureWarning: `distplot` is a deprecated function and will be removed
```

in a future version. Please adapt your code to use either ``displot`` (a figure-level function with similar flexibility) or ``histplot`` (an axes-level function for histograms).


```
warnings.warn(msg, FutureWarning)
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7ff47af787d0>
```



```
sns.histplot(data['Age'])
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7ff477ed9d10>
```




```
sns.boxplot(data['Age'])
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43:  
FutureWarning: Pass the following variable as a keyword arg: x. From  
version 0.12, the only valid positional argument will be `data`, and  
passing other arguments without an explicit keyword will result in an  
error or misinterpretation.
```


```
FutureWarning
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7ff477d9ea90>
```




3.2 Bi-variate Analysis


```
plt.figure(figsize=(7,7))  
sns.lineplot(data = data, x = 'Tenure', y = 'CreditScore')  
<matplotlib.axes._subplots.AxesSubplot at 0x7ff47afc7750>
```



```
plt.figure(figsize=(10,10))
sns.barplot(data = data, x = 'CreditScore', y = 'EstimatedSalary')
<matplotlib.axes._subplots.AxesSubplot at 0x7ff477da5e10>
```

```
plt.figure(figsize=(10,10))
sns.barplot(data = data, x = 'CreditScore', y = 'Tenure')
<matplotlib.axes._subplots.AxesSubplot at 0x7ff47b1b3cd0>
```




```
plt.figure(figsize=(10,10))
sns.lineplot(data['Age'], data['EstimatedSalary'])
```

/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43:
FutureWarning: Pass the following variables as keyword args: x, y.
From version 0.12, the only valid positional argument will be `data`,
and passing other arguments without an explicit keyword will result in
an error or misinterpretation.

FutureWarning

<matplotlib.axes._subplots.AxesSubplot at 0x7ff475933690>




```
plt.figure(figsize=(17,17))
sns.barplot(data['Age'], data['EstimatedSalary'])
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43:
FutureWarning: Pass the following variables as keyword args: x, y.
From version 0.12, the only valid positional argument will be `data`,
and passing other arguments without an explicit keyword will result in
an error or misinterpretation.
```

```
FutureWarning
```


```
<matplotlib.axes._subplots.AxesSubplot at 0x7ff4749085d0>
```



3.3 Multivariate Analysis


```
sns.scatterplot(data = data, x = 'CreditScore', y = 'Age')
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7ff472f36a10>
```



```
sns.scatterplot(data = data, x = 'CreditScore', y = 'Balance', hue = 'Gender')
```

<matplotlib.axes._subplots.AxesSubplot at 0x7ff472e6d3d0>




```
sns.scatterplot(data['Tenure'], data['CreditScore'], hue = data['Gender'])
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43:
FutureWarning: Pass the following variables as keyword args: x, y.
From version 0.12, the only valid positional argument will be `data`,
and passing other arguments without an explicit keyword will result in
an error or misinterpretation.
```

FutureWarning

<matplotlib.axes._subplots.AxesSubplot at 0x7ff472ef03d0>




```
sns.scatterplot(data['Age'], data['Balance'], hue = data['Gender'])
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43:
FutureWarning: Pass the following variables as keyword args: x, y.
From version 0.12, the only valid positional argument will be `data`,
and passing other arguments without an explicit keyword will result in
an error or misinterpretation.
```

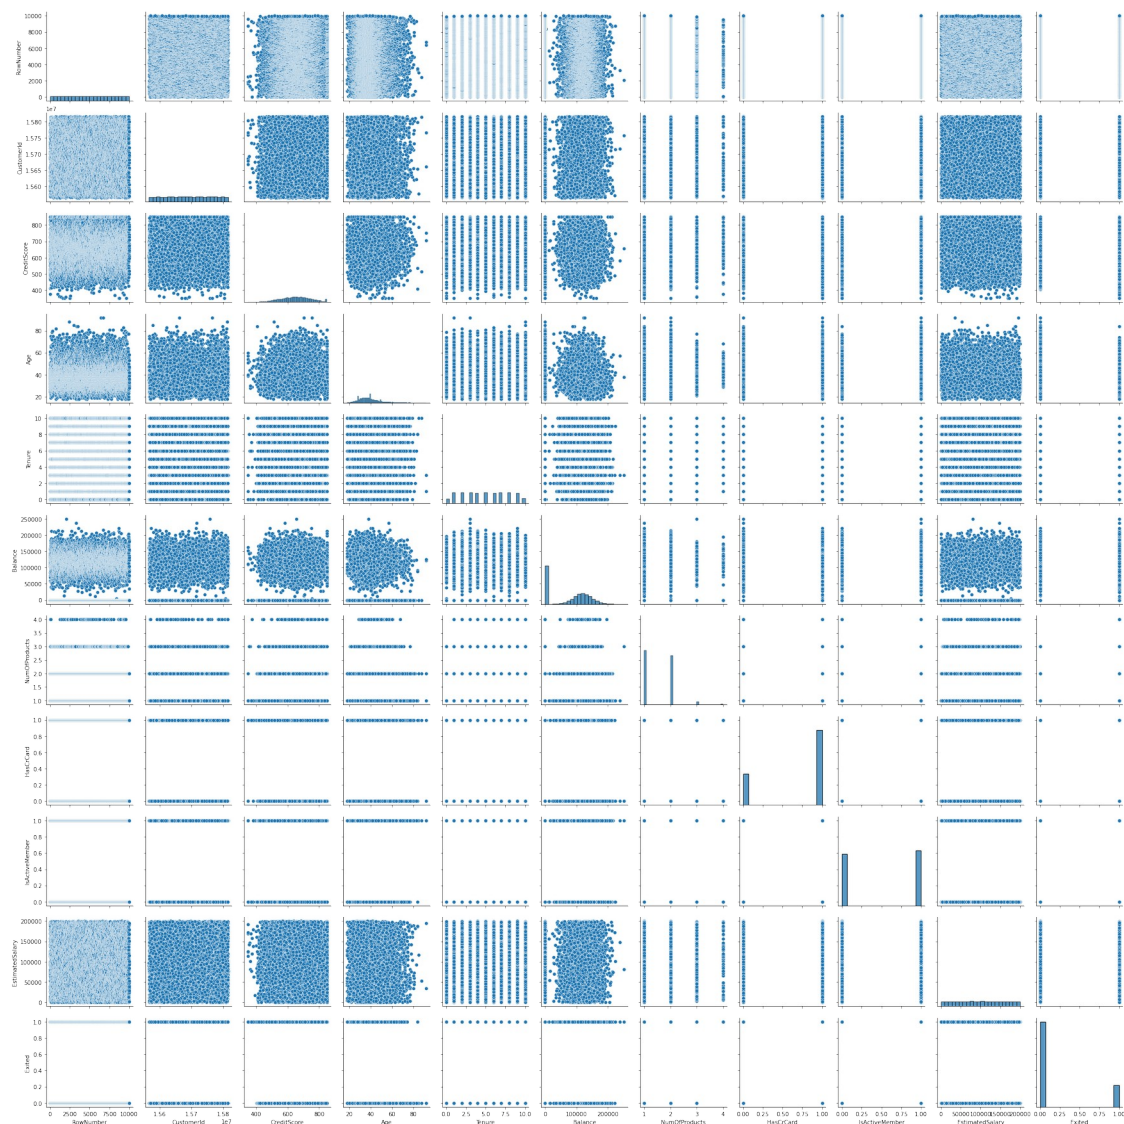
FutureWarning

<matplotlib.axes._subplots.AxesSubplot at 0x7ff472dacf10>



```
sns.pairplot(data)
```

```
<seaborn.axisgrid.PairGrid at 0x7ff472db8650>
```

Question-4. Perform descriptive statistics on the dataset

```
data.mean(numeric_only = True)
```

```
RowNumber      5.000500e+03
CustomerId     1.569094e+07
CreditScore    6.505288e+02
Age            3.892180e+01
Tenure         5.012800e+00
Balance        7.648589e+04
NumOfProducts  1.530200e+00
HasCrCard      7.055000e-01
IsActiveMember 5.151000e-01
EstimatedSalary 1.000902e+05
Exited         2.037000e-01
dtype: float64
```



```
data.median(numeric_only = True)
```

```
RowNumber      5.000500e+03
CustomerId      1.569074e+07
CreditScore     6.520000e+02
Age             3.700000e+01
Tenure          5.000000e+00
Balance         9.719854e+04
NumOfProducts   1.000000e+00
HasCrCard       1.000000e+00
IsActiveMember  1.000000e+00
EstimatedSalary 1.001939e+05
Exited          0.000000e+00
dtype: float64
```

```
data['CreditScore'].mode()
```

```
0      850
dtype: int64
```

```
data['EstimatedSalary'].mode()
```

```
0      24924.92
dtype: float64
```

```
data['HasCrCard'].unique()
```

```
array([1, 0])
```

```
data['Tenure'].unique()
```

```
array([ 2,  1,  8,  7,  4,  6,  3, 10,  5,  9,  0])
```

```
data.std(numeric_only=True)
```

```
RowNumber      2886.895680
CustomerId      71936.186123
CreditScore     96.653299
Age             10.487806
Tenure          2.892174
Balance         62397.405202
NumOfProducts    0.581654
HasCrCard        0.455840
IsActiveMember   0.499797
EstimatedSalary 57510.492818
Exited          0.402769
dtype: float64
```

```
data.describe()
```

	RowNumber	CustomerId	CreditScore	Age
Tenure \				
count	10000.000000	1.000000e+04	10000.000000	10000.000000

```

10000.000000
mean      5000.50000  1.569094e+07    650.528800    38.921800
5.012800
std       2886.89568  7.193619e+04     96.653299    10.487806
2.892174
min        1.00000  1.556570e+07    350.000000    18.000000
0.000000
25%       2500.75000  1.562853e+07    584.000000    32.000000
3.000000
50%       5000.50000  1.569074e+07    652.000000    37.000000
5.000000
75%       7500.25000  1.575323e+07    718.000000    44.000000
7.000000
max      10000.00000  1.581569e+07    850.000000    92.000000
10.000000

```

	Balance	NumOfProducts	HasCrCard	IsActiveMember \
count	10000.000000	10000.000000	10000.00000	10000.000000
mean	76485.889288	1.530200	0.70550	0.515100
std	62397.405202	0.581654	0.45584	0.499797
min	0.000000	1.000000	0.00000	0.000000
25%	0.000000	1.000000	0.00000	0.000000
50%	97198.540000	1.000000	1.00000	1.000000
75%	127644.240000	2.000000	1.00000	1.000000
max	250898.090000	4.000000	1.00000	1.000000

	EstimatedSalary	Exited
count	10000.000000	10000.000000
mean	100090.239881	0.203700
std	57510.492818	0.402769
min	11.580000	0.000000
25%	51002.110000	0.000000
50%	100193.915000	0.000000
75%	149388.247500	0.000000
max	199992.480000	1.000000

Question-5. Handle the Missing values.

```
data.isnull().any()
```

RowNumber	False
CustomerId	False
Surname	False
CreditScore	False
Geography	False
Gender	False
Age	False
Tenure	False
Balance	False
NumOfProducts	False
HasCrCard	False

```
IsActiveMember    False
EstimatedSalary   False
Exited            False
dtype: bool
```

```
data.isnull().sum()
```

```
RowNumber          0
CustomerId          0
Surname            0
CreditScore        0
Geography          0
Gender             0
Age               0
Tenure            0
Balance            0
NumOfProducts     0
HasCrCard          0
IsActiveMember     0
EstimatedSalary    0
Exited             0
dtype: int64
```

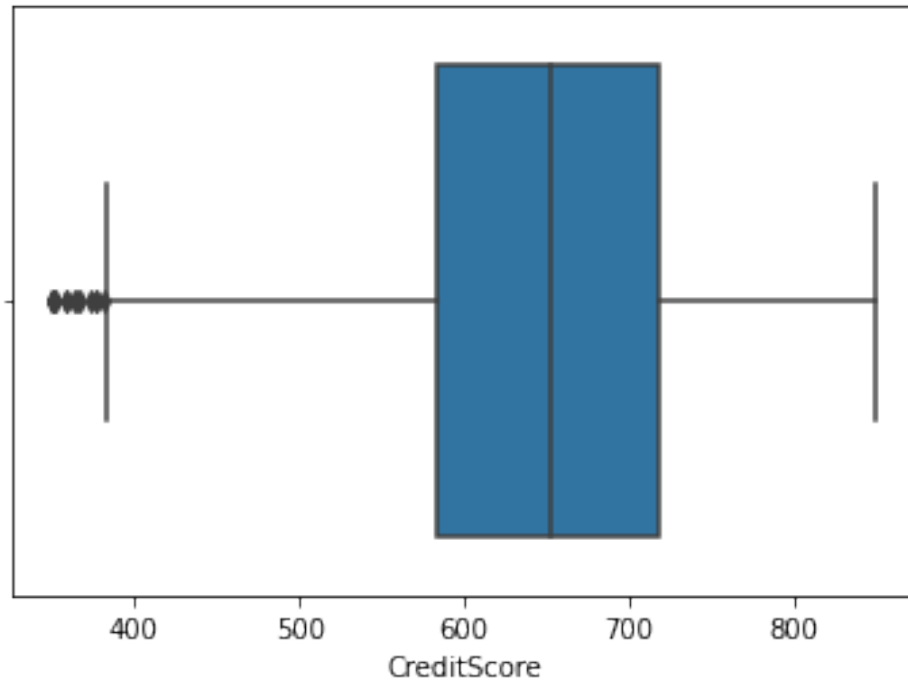
Question-6. Find the outliers and replace the outliers

```
sns.boxplot(data['CreditScore']) #Outlier detection - box plot
```


```
/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43:
FutureWarning: Pass the following variable as a keyword arg: x. From
version 0.12, the only valid positional argument will be `data`, and
passing other arguments without an explicit keyword will result in an
error or misinterpretation.
```

```
FutureWarning
```


```
<matplotlib.axes._subplots.AxesSubplot at 0x7ff46fbdf3d0>
```



```
fig, ax = plt.subplots(figsize = (5,3)) #Outlier detection - Scatter
plot
ax.scatter(data['Balance'], data['Exited'])
# x-axis label
ax.set_xlabel('Balance')
# y-axis label
ax.set_ylabel('Exited')
plt.show()
sns.boxplot(x=data['Balance'])
```



<matplotlib.axes._subplots.AxesSubplot at 0x7ff47177b710>



```
from scipy import stats #Outlier detection – zscore
zscore = np.abs(stats.zscore(data['CreditScore']))
print(zscore)
print('No. of Outliers : ', np.shape(np.where(zscore>3)))
```

```
0      0.326221
1      0.440036
2      1.536794
3      0.501521
4      2.063884
```

```
...
9995   1.246488
9996   1.391939
9997   0.604988
9998   1.256835
9999   1.463771
```

Name: CreditScore, Length: 10000, dtype: float64

No. of Outliers : (1, 8)

```
q = data.quantile([0.75,0.25])
```

q

	RowNumber	CustomerId	CreditScore	Age	Tenure	Balance	\
0.75	7500.25	15753233.75	718.0	44.0	7.0	127644.24	
0.25	2500.75	15628528.25	584.0	32.0	3.0	0.00	

	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary
Exited				
0.75	2.0	1.0	1.0	149388.2475

```
0.0
0.25          1.0          0.0          0.0          51002.1100
0.0
```

```
iqr = q.iloc[0] - q.iloc[1]
iqr
```

```
RowNumber      4999.5000
CustomerId     124705.5000
CreditScore    134.0000
Age            12.0000
Tenure         4.0000
Balance       127644.2400
NumOfProducts  1.0000
HasCrCard      1.0000
IsActiveMember 1.0000
EstimatedSalary 98386.1375
Exited         0.0000
dtype: float64
```

```
u = q.iloc[0] + (1.5*iqr)
u
```

```
RowNumber      1.499950e+04
CustomerId     1.594029e+07
CreditScore    9.190000e+02
Age            6.200000e+01
Tenure         1.300000e+01
Balance       3.191106e+05
NumOfProducts  3.500000e+00
HasCrCard      2.500000e+00
IsActiveMember 2.500000e+00
EstimatedSalary 2.969675e+05
Exited         0.000000e+00
dtype: float64
```

```
l = q.iloc[1] - (1.5*iqr)
l
```

```
RowNumber      -4.998500e+03
CustomerId     1.544147e+07
CreditScore    3.830000e+02
Age            1.400000e+01
Tenure         -3.000000e+00
Balance       -1.914664e+05
NumOfProducts  -5.000000e-01
HasCrCard      -1.500000e+00
IsActiveMember -1.500000e+00
EstimatedSalary -9.657710e+04
Exited         0.000000e+00
dtype: float64
```

```

Q1 = data['EstimatedSalary'].quantile(0.25) #Outlier detection - IQR
Q3 = data['EstimatedSalary'].quantile(0.75)
iqr = Q3 - Q1
print(iqr)
upper=Q3 + 1.5 * iqr
lower=Q1 - 1.5 * iqr
count = np.size(np.where(data['EstimatedSalary'] >upper))
count = count + np.size(np.where(data['EstimatedSalary'] <lower))
print('No. of outliers : ', count)

```

98386.1375

No. of outliers : 0

```


data['CreditScore'] = np.where(np.logical_or(data['CreditScore']>900,
data['CreditScore']<383), 650, data['CreditScore'])
sns.boxplot(data['CreditScore'])

```

/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43:
FutureWarning: Pass the following variable as a keyword arg: x. From
version 0.12, the only valid positional argument will be `data`, and
passing other arguments without an explicit keyword will result in an
error or misinterpretation.

FutureWarning

<matplotlib.axes._subplots.AxesSubplot at 0x7ff47b1b5910>



```

upper = data.Age.mean() + (3 * data.Age.std()) #Outlier detection - 3
sigma
lower = data.Age.mean() - (3 * data.Age.std())
columns = data[ ( data['Age'] > upper ) | ( data['Age']<lower ) ]

```

```

print('Upper range : ', upper)
print('Lower range : ', lower)
print('No. of Outliers : ', len(columns))

Upper range : 70.38521935511383
Lower range : 7.458380644886169
No. of Outliers : 133

columns = ['EstimatedSalary', 'Age', 'Balance', 'NumOfProducts',
'Tenure', 'CreditScore']
#After outlier removal
for i in columns:
    Q1 = data[i].quantile(0.25)
    Q3 = data[i].quantile(0.75)
    iqr = Q3 - Q1
    upper=Q3 + 1.5 * iqr
    lower=Q1 - 1.5 * iqr
    count = np.size(np.where(data[i] >upper))
    count = count + np.size(np.where(data[i] <lower))
    print('No. of outliers in ', i, ' : ', count)

```

```

No. of outliers in EstimatedSalary : 0
No. of outliers in Age : 359
No. of outliers in Balance : 0
No. of outliers in NumOfProducts : 60
No. of outliers in Tenure : 0
No. of outliers in CreditScore : 0

```

Question-6. Find the outliers and replace the outliers

```

from sklearn.preprocessing import LabelEncoder, OneHotEncoder
le = LabelEncoder()
oneh = OneHotEncoder()
data['Surname'] = le.fit_transform(data['Surname'])
data['Gender'] = le.fit_transform(data['Gender'])
data['Geography'] = le.fit_transform(data['Geography'])
data.head()

```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age
0	1	15634602	1115	619	0	0	42
1	2	15647311	1177	608	2	0	41
2	3	15619304	2040	502	0	0	42
3	4	15701354	289	699	0	0	39
4	5	15737888	1822	850	2	0	43

	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	\
0	2	0.00	1	1	1	
1	1	83807.86	1	0	1	
2	8	159660.80	3	1	0	
3	1	0.00	2	0	0	
4	2	125510.82	1	1	1	

	EstimatedSalary	Exited
0	101348.88	1
1	112542.58	0
2	113931.57	1
3	93826.63	0
4	79084.10	0

Question-8. Split the data into dependent and independent variables split the data in X and Y

```
x = data.iloc[:, 0:13]
x # independent values ( inputs)
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender
Age \						
0	1	15634602	1115	619	0	0
42						
1	2	15647311	1177	608	2	0
41						
2	3	15619304	2040	502	0	0
42						
3	4	15701354	289	699	0	0
39						
4	5	15737888	1822	850	2	0
43						
...
...						
9995	9996	15606229	1999	771	0	1
39						
9996	9997	15569892	1336	516	0	1
35						
9997	9998	15584532	1570	709	0	0
36						
9998	9999	15682355	2345	772	1	1
42						
9999	10000	15628319	2751	792	0	0
28						

	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	\
0	2	0.00	1	1	1	
1	1	83807.86	1	0	1	
2	8	159660.80	3	1	0	
3	1	0.00	2	0	0	

4	2	125510.82	1	1	1
...
9995	5	0.00	2	1	0
9996	10	57369.61	1	1	1
9997	7	0.00	1	0	1
9998	3	75075.31	2	1	0
9999	4	130142.79	1	1	0

	EstimatedSalary
0	101348.88
1	112542.58
2	113931.57
3	93826.63
4	79084.10
...	...
9995	96270.64
9996	101699.77
9997	42085.58
9998	92888.52
9999	38190.78

[10000 rows x 13 columns]

```
y = data['Exited']
y # dependent values (output)
```

0	1
1	0
2	1
3	0
4	0
...	...
9995	0
9996	0
9997	1
9998	1
9999	0

Name: Exited, Length: 10000, dtype: int64

Question-9. Scale the independent variables

```
from sklearn.preprocessing import StandardScaler, MinMaxScaler
sc = StandardScaler()
x_scaled = sc.fit_transform(x)
x_scaled
```

array([[-1.73187761, -0.78321342, -0.46418322, ..., 0.64609167,
 0.97024255, 0.02188649],
 [-1.7315312 , -0.60653412, -0.3909112 , ..., -1.54776799,
 0.97024255, 0.21653375],
 [-1.73118479, -0.99588476, 0.62898807, ..., 0.64609167,

```

-1.03067011, 0.2406869 ],
...,
[ 1.73118479, -1.47928179, 0.07353887, ..., -1.54776799,
  0.97024255, -1.00864308],
[ 1.7315312 , -0.11935577, 0.98943914, ..., 0.64609167,
  -1.03067011, -0.12523071],
[ 1.73187761, -0.87055909, 1.4692527 , ..., 0.64609167,
  -1.03067011, -1.07636976]])

```

Question-10. Split x and y into Training and Testing

```

from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x_scaled, y,
test_size = 0.3, random_state = 0)

```

x_train

```

array([[ 0.92889885, -0.79703192, -1.47580983, ..., 0.64609167,
         0.97024255, -0.77021814],
       [ 1.39655257, 0.71431365, -1.58808148, ..., 0.64609167,
        -1.03067011, -1.39576675],
       [-0.4532777 , 0.96344969, -0.24082173, ..., -1.54776799,
         0.97024255, -1.49965629],
       ...,
       [-0.60119484, -1.62052514, -0.36136603, ..., 0.64609167,
        -1.03067011, 1.41441489],
       [ 1.67853045, -0.37403866, 0.72589622, ..., 0.64609167,
         0.97024255, 0.84614739],
       [-0.78548505, -1.36411841, 1.3829808 , ..., 0.64609167,
        -1.03067011, 0.32630495]])

```

x_train.shape

```
(7000, 13)
```

x_test

```

array([[ 1.52229946, -1.04525042, 1.39834429, ..., 0.64609167,
         0.97024255, 1.61304597],
       [-1.42080128, -0.50381294, -0.78208925, ..., 0.64609167,
        -1.03067011, 0.49753166],
       [-0.90118604, -0.7932923 , 0.41271742, ..., 0.64609167,
         0.97024255, -0.4235611 ],
       ...,
       [ 1.49216178, -0.14646448, 0.6868966 , ..., 0.64609167,
         0.97024255, 1.17045451],
       [ 1.1758893 , -1.29228727, -1.38481071, ..., 0.64609167,
         0.97024255, -0.50846777],
       [ 0.08088677, -1.38538833, 1.11707427, ..., 0.64609167,
         0.97024255, -1.15342685]])

```

x_test.shape

```
(3000, 13)
```

```
y_train
```

```
7681    1
9031    0
3691    0
202     1
5625    0
```

```
..
9225    0
4859    0
3264    0
9845    0
2732    1
```

```
Name: Exited, Length: 7000, dtype: int64
```

```
y_test
```

```
9394    0
898     1
2398    0
5906    0
2343    0
```

```
..
4004    0
7375    0
9307    0
8394    0
5233    1
```

```
Name: Exited, Length: 3000, dtype: int64
```